



AI + DOCUMENTS

Docu-Genius: AI-Powered Document System

Demo Overview

An end-to-end, scalable, AI-driven document insight platform

 Presenter: [Your Name] |  GitHub: lihuiniu/docu-genius



Demo Agenda

- 1 What is Docu-Genius?
- 2 Key Features
- 3 Architecture Overview
- 4 Model Choices: **OpenAI** + **Seq2Seq** + **LangChain**
- 5 Workflow & Controller: **LangGraph**
- 6 Technical Details & API Usage
- 7 Demo Highlights
- 8 Q&A



What is Docu-Genius?

AI-powered system for document chunking, semantic search, summarization, and evaluation



FastAPI-powered async backend

High-performance, scalable server with asynchronous processing



LLM-powered summarization & evaluation

Advanced AI for content generation and quality control



Semantic retrieval & caching

Milvus vector search with Redis caching for fast access



Flexible storage & interfaces

Local & cloud storage support with CLI & API access



Key Features



Flexible Storage

Local, S3, Azure, Delta Lake



Semantic Search

Milvus vector database



Fast Cache

Redis 8.0 for quick retrieval



AI Evaluation

OpenAI LLMs (v1.x), LangChain LLMChain



Workflow Orchestration

LangGraph StateGraph workflows



Async & Scalable

FastAPI + Hypercorn server

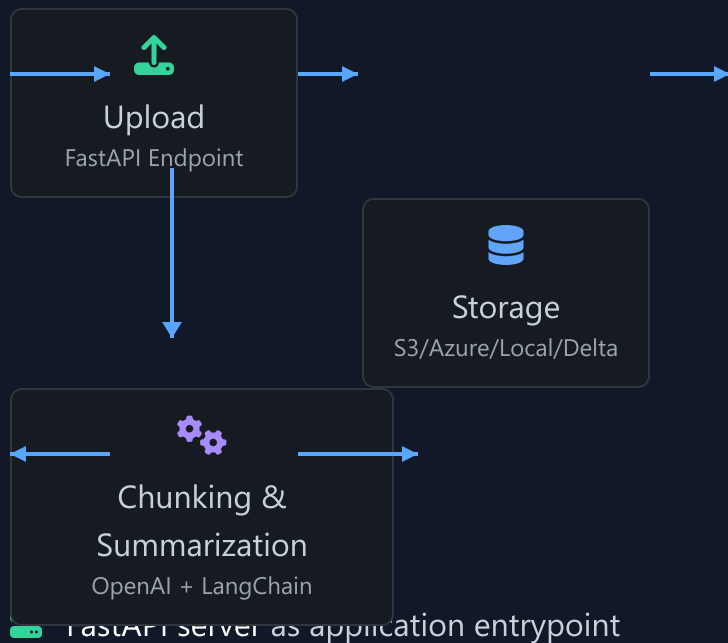


CLI Utility

Batch/parallel reindexing with resume support



System Architecture



FastAPI server as application endpoint

OpenAI + LangChain for embeddings & summarization

CLI utility for batch operations & admin tasks

LangGraph for stateful workflow orchestration

Milvus + Redis for semantic retrieval & caching

Async pipeline for scalable processing



Model Choices & Pipeline

Integrated AI components for document understanding and evaluation

OpenAI

- Embeddings for semantic search
- Document summarization
- OpenAI API v1.x integration

Seq2Seq

- Text transformation approach
- Via OpenAI's API
- Abstracted through LangChain

LangChain

- Workflow integration
- LLMChain for evaluation
- Modular component chaining

Quality Control Pipeline

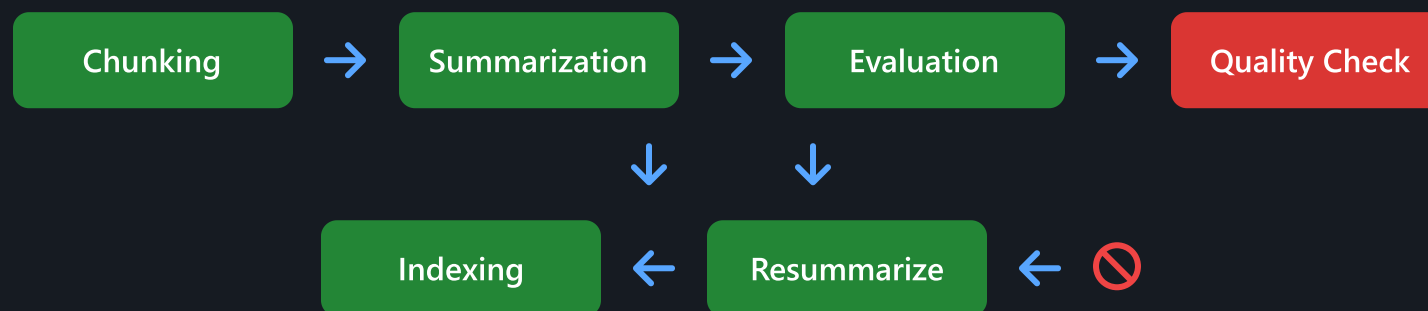




Workflow Orchestration: LangGraph

StateGraph-based workflow for complex document processing pipelines

LangGraph StateGraph Workflow



Modular & Conditional Flows

Dynamic routing based on quality evaluations



Unified Orchestration

Same workflow powers CLI batch jobs & API requests



Async Execution

Non-blocking processing for high throughput



Extensible Architecture

Easy to add new states or decision logic



Technical Implementation & Demo

⚙️ Setup

- Python virtual environment + requirements.txt
- FastAPI + Hypercorn for async API
- Milvus 2.6 for vector storage
- Redis 8.0 for caching
- OpenAI API key for LLM functionality

🔧 Testing

- Unit & integration tests with pytest
- Async testing with pytest-asyncio
- Continuous integration via GitHub Actions
- Test coverage reporting

</> API Usage Examples

Upload document:

```
curl -F "file=@./example.txt" http://localhost:8000/upload/local
```

Summarize document:

```
curl -X POST http://localhost:8000/summarize \
-H "Content-Type: application/json" \
-d '{"doc_id": "your_doc_id", "storage": "local"}'
```

Query similar chunks:

```
curl -X POST http://localhost:8000/query \
-H "Content-Type: application/json" \
-d '{"keyword": "example", "top_k": 5}'
```

Batch reindex via CLI:

```
python cli/reindex.py --doc-ids doc1 doc2 --storage local --concurrency
```

💡 Perfect for real-time document intelligence in modern applications